

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('winequality-red.csv')
df
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0

1599 rows × 12 columns



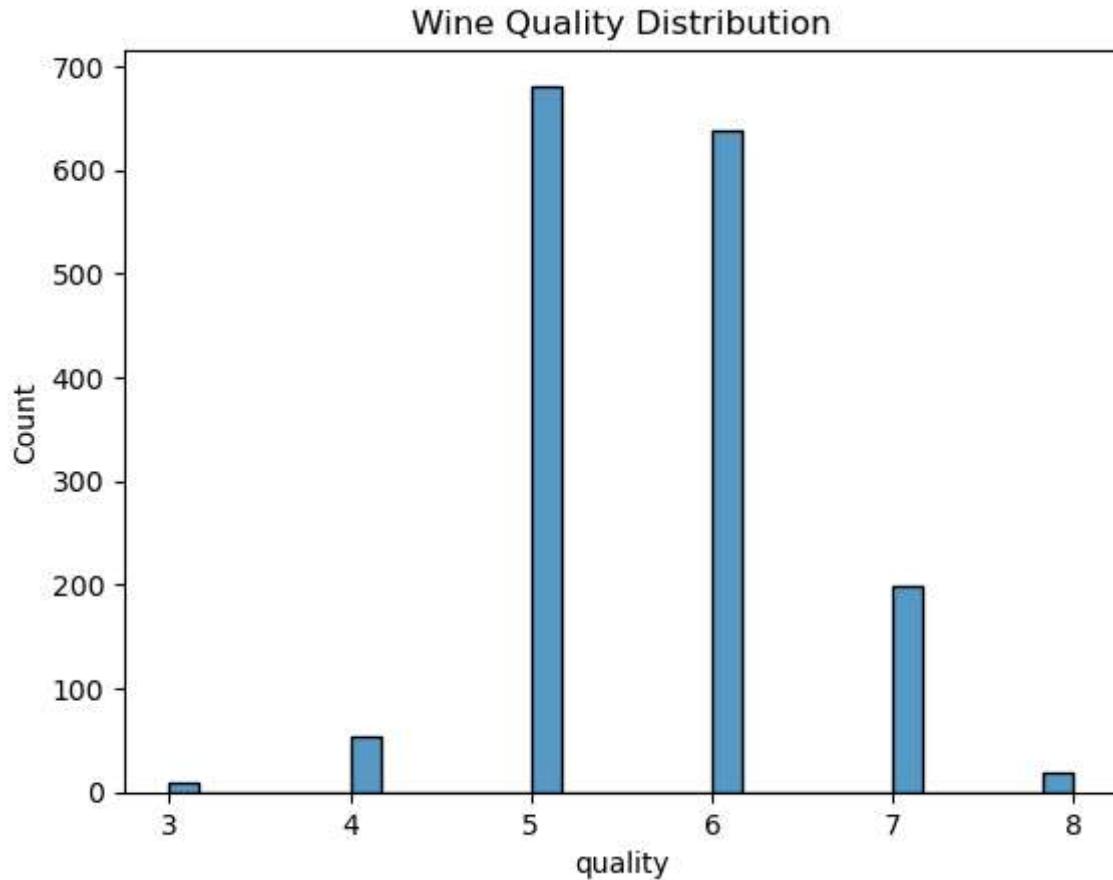
```
In [3]: df.isnull().sum()
```

```
Out[3]: fixed acidity      0
volatile acidity      0
citric acid          0
residual sugar        0
chlorides            0
free sulfur dioxide  0
total sulfur dioxide 0
density              0
pH                   0
sulphates            0
alcohol              0
quality              0
dtype: int64
```

```
In [4]: display(list(df))
sns.histplot(x=df.quality)
plt.title("Wine Quality Distribution")
```

```
['fixed acidity',
 'volatile acidity',
 'citric acid',
 'residual sugar',
 'chlorides',
 'free sulfur dioxide',
 'total sulfur dioxide',
 'density',
 'pH',
 'sulphates',
 'alcohol',
 'quality']
```

Out[4]: Text(0.5, 1.0, 'Wine Quality Distribution')



In [5]: df.describe()

## RED WINE

Out[5]:	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	
<b>count</b>	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	15
<b>mean</b>	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	
<b>std</b>	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	
<b>min</b>	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	
<b>25%</b>	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	
<b>50%</b>	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	
<b>75%</b>	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	
<b>max</b>	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	

◀ ▶

In [6]: `df.shape`

Out[6]: `(1599, 12)`

In [7]: `df.duplicated().sum()`

Out[7]: `240`

In [8]: `df.drop_duplicates(keep = 'first', inplace = True, ignore_index = True)`  
`df.shape`

Out[8]: `(1359, 12)`

In [9]: `df['quality'].value_counts()`

Out[9]: `quality`  
5 577  
6 535  
7 167  
4 53  
8 17  
3 10  
Name: count, dtype: int64

In [10]: `df['Class'] = np.where(df['quality'] >= 7, 1, 0)`

In [11]: `df`

## RED WINE

Out[11]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
4	7.4	0.660	0.00	1.8	0.075	13.0	40.0	0.99780	3.51	0.56	9.4
...	...	...	...	...	...	...	...	...	...	...	...
1354	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82	9.5
1355	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5
1356	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2
1357	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2
1358	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0

1359 rows × 13 columns

In [12]: `plt.figure(figsize = (15,12))  
sns.heatmap(df.corr(), annot = True, cmap = 'Pastel1')  
plt.show()`

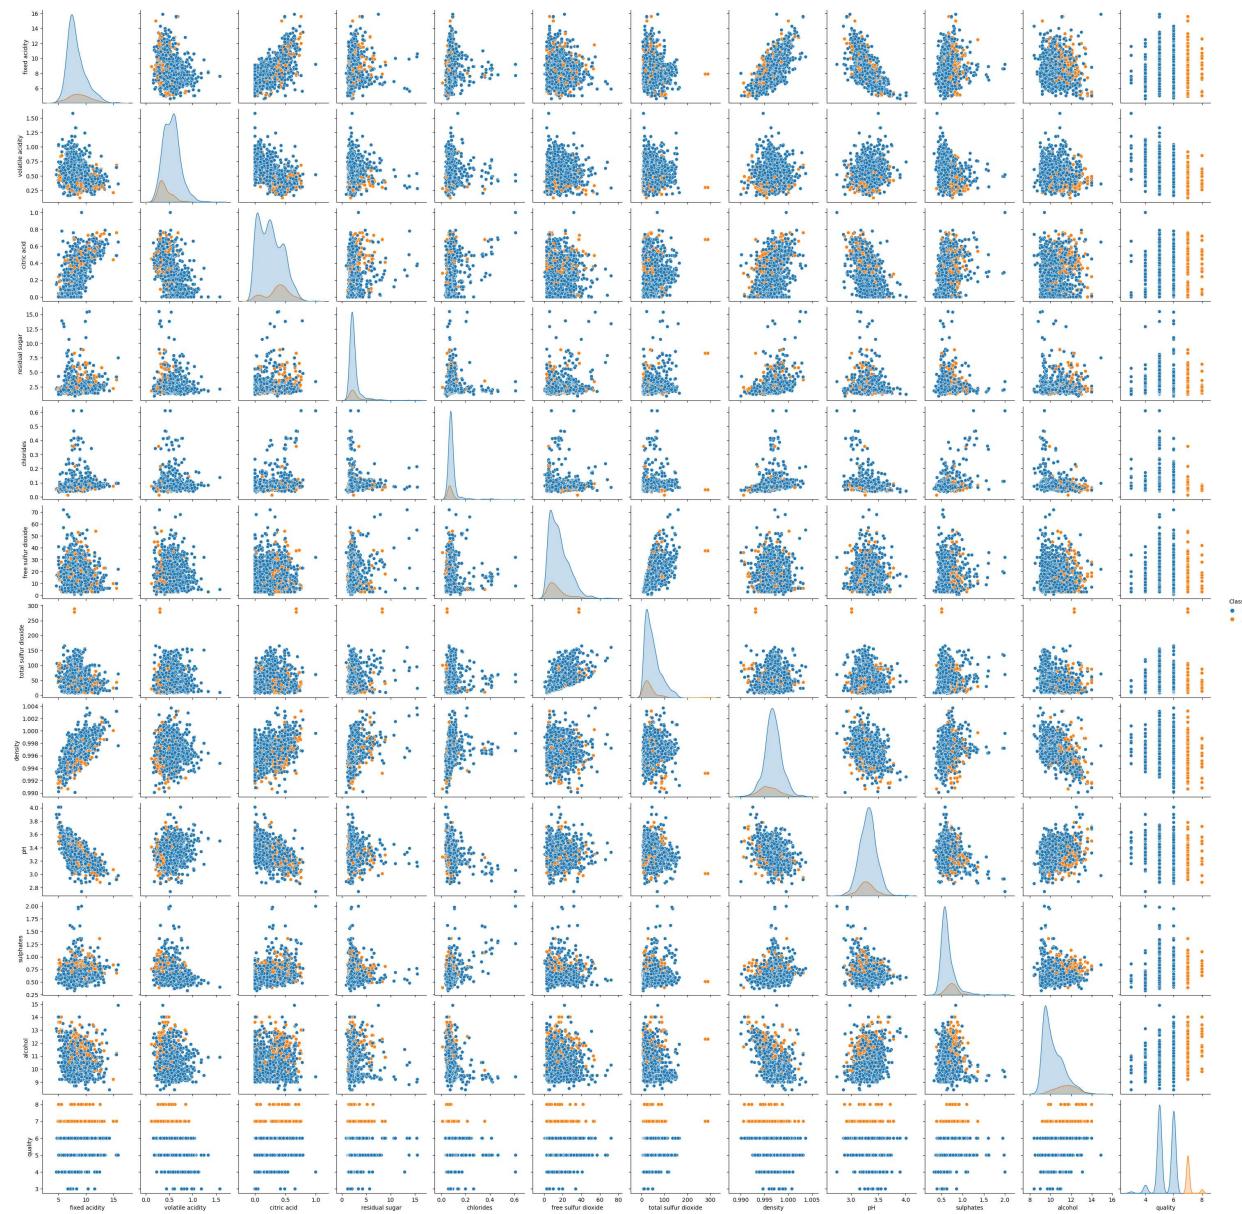
## RED WINE



```
In [13]: plt.figure(figsize = (15,20))
sns.pairplot(data=df, hue='Class')
plt.show()
```

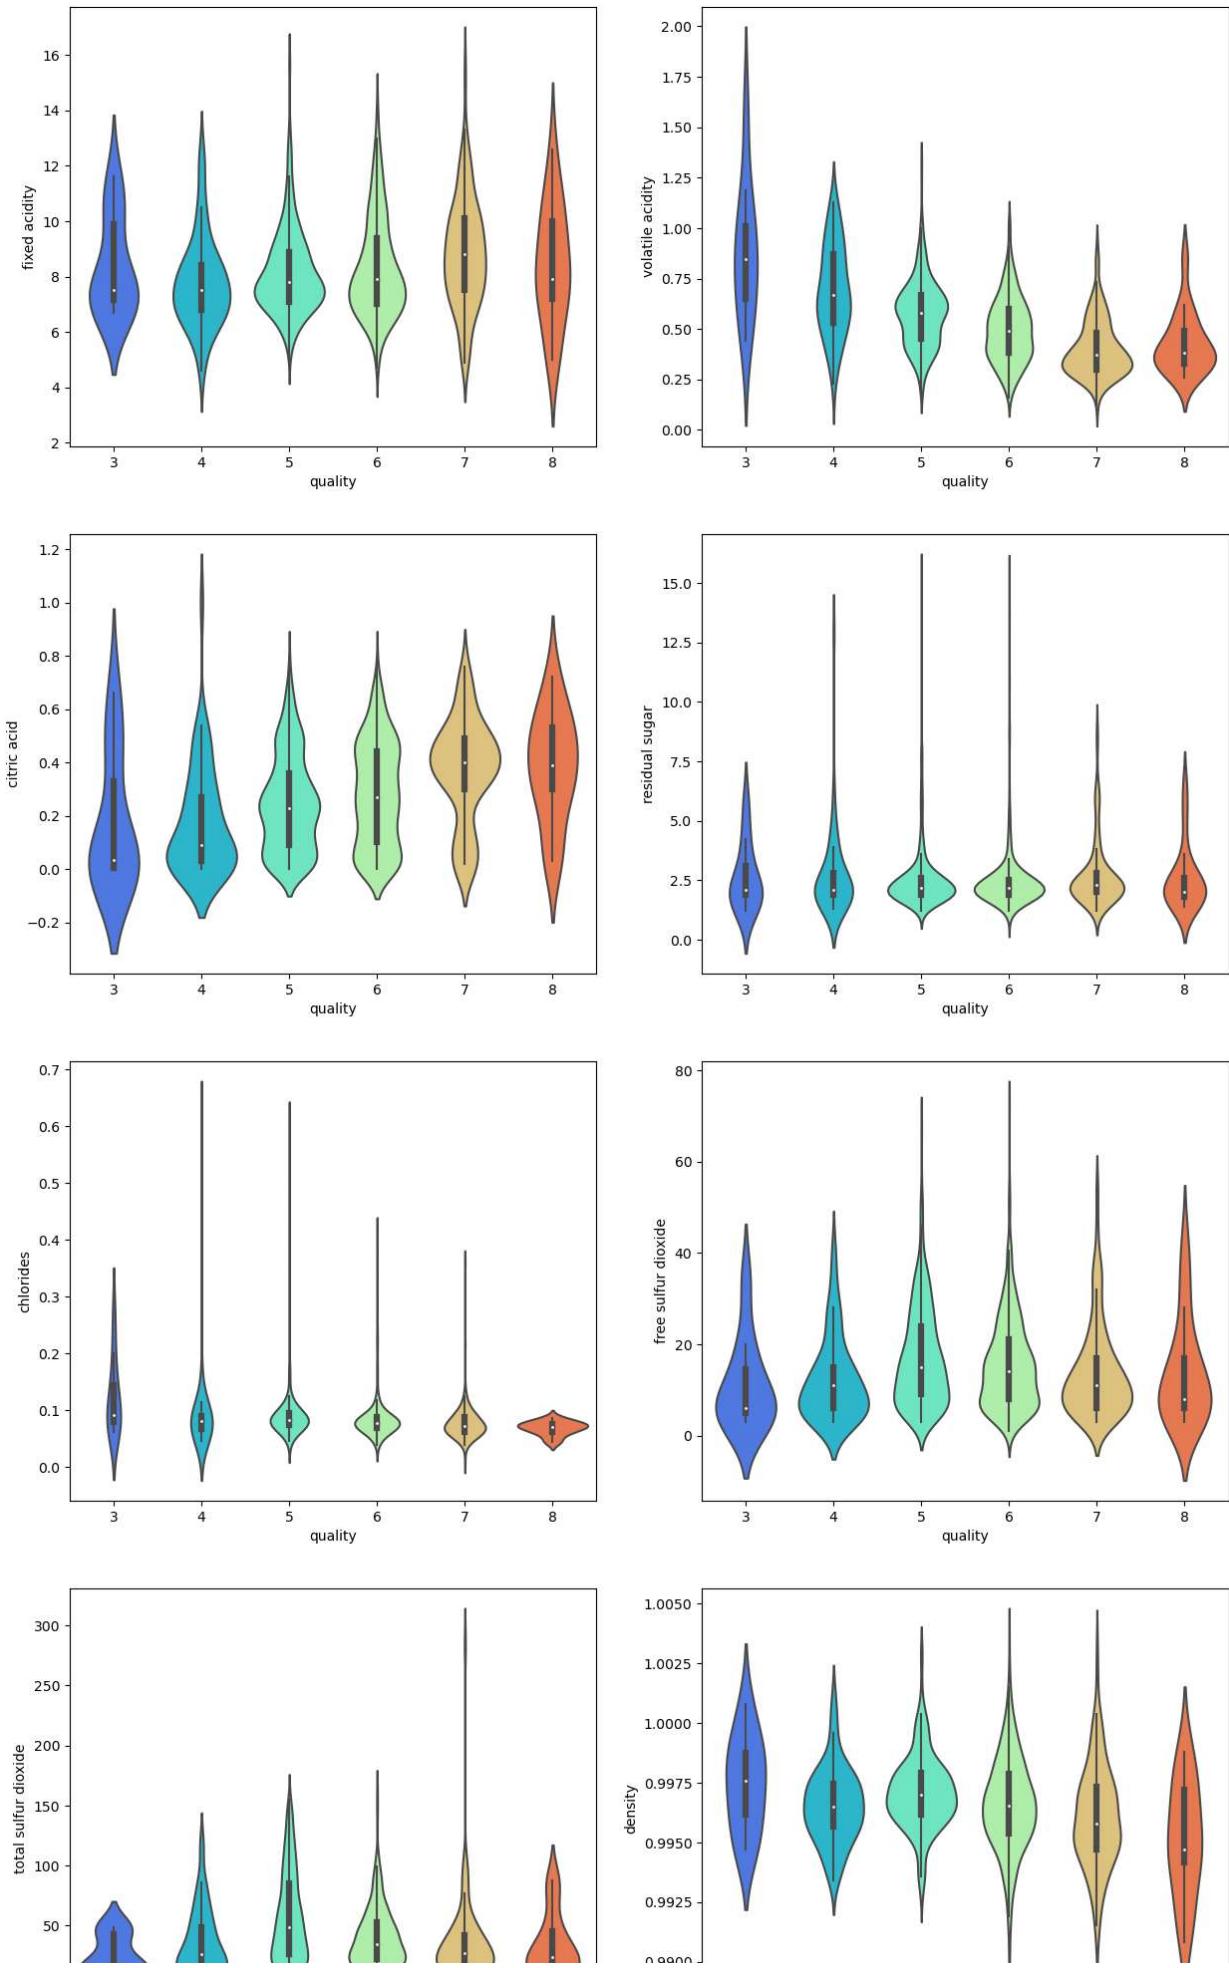
<Figure size 1500x2000 with 0 Axes>

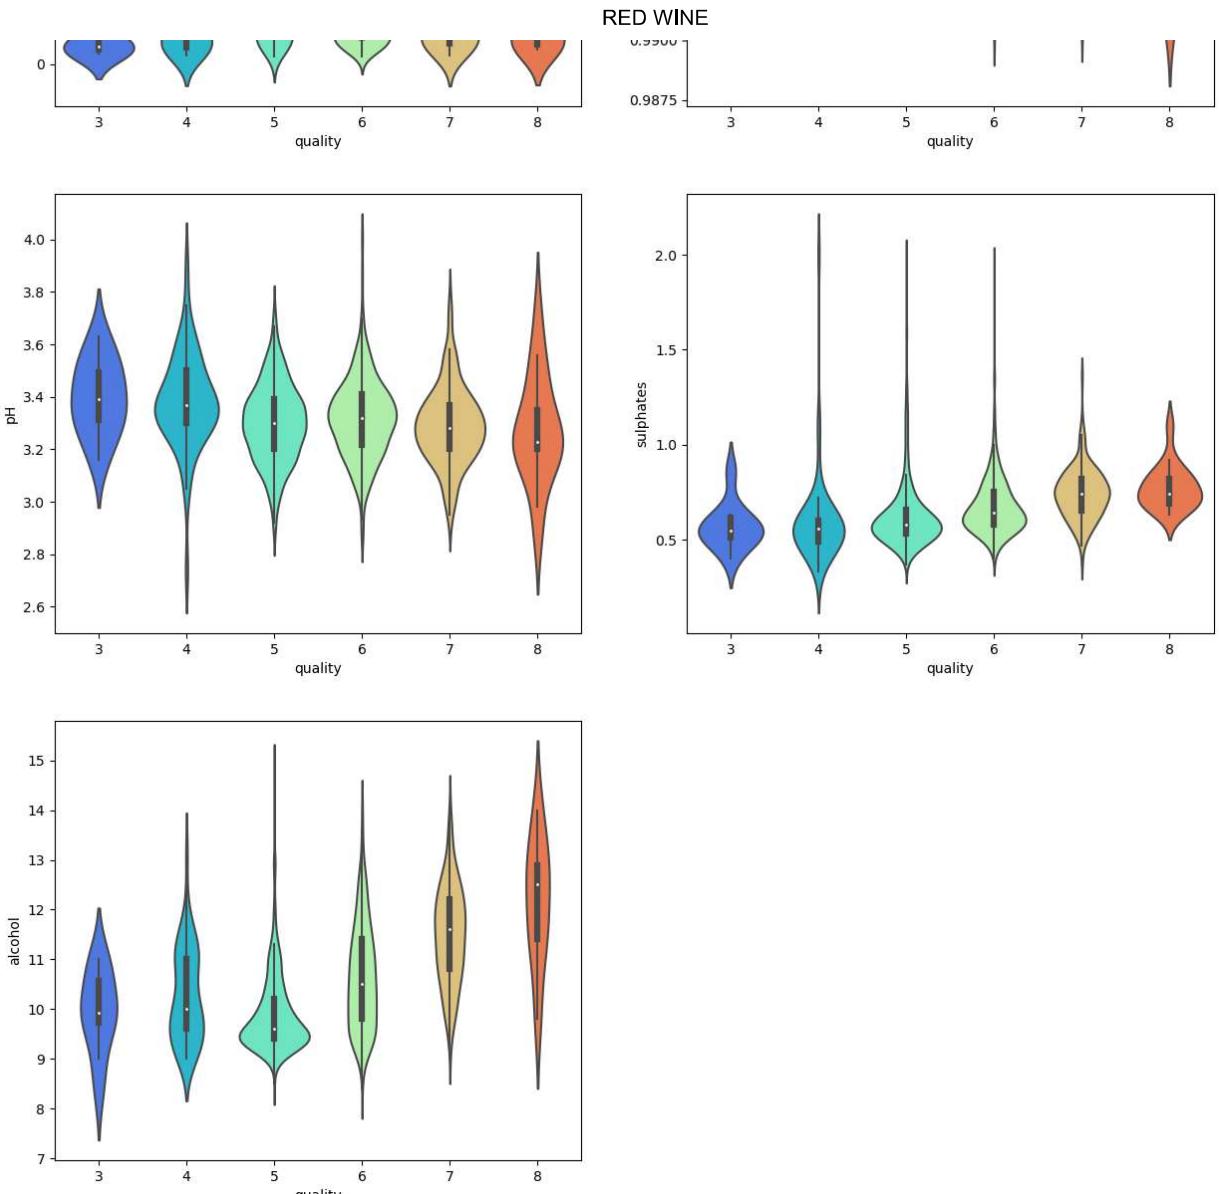
## RED WINE



```
In [14]: plt.figure(figsize=(15,150))
plot_number = 1
for i in list(df.select_dtypes(include=['float']).columns)[0:]:
    plt.subplot(22,2,plot_number)
    sns.violinplot(x="quality", y = i, data=df, palette='rainbow')
    plot_number+=1
plt.show()
```

## RED WINE





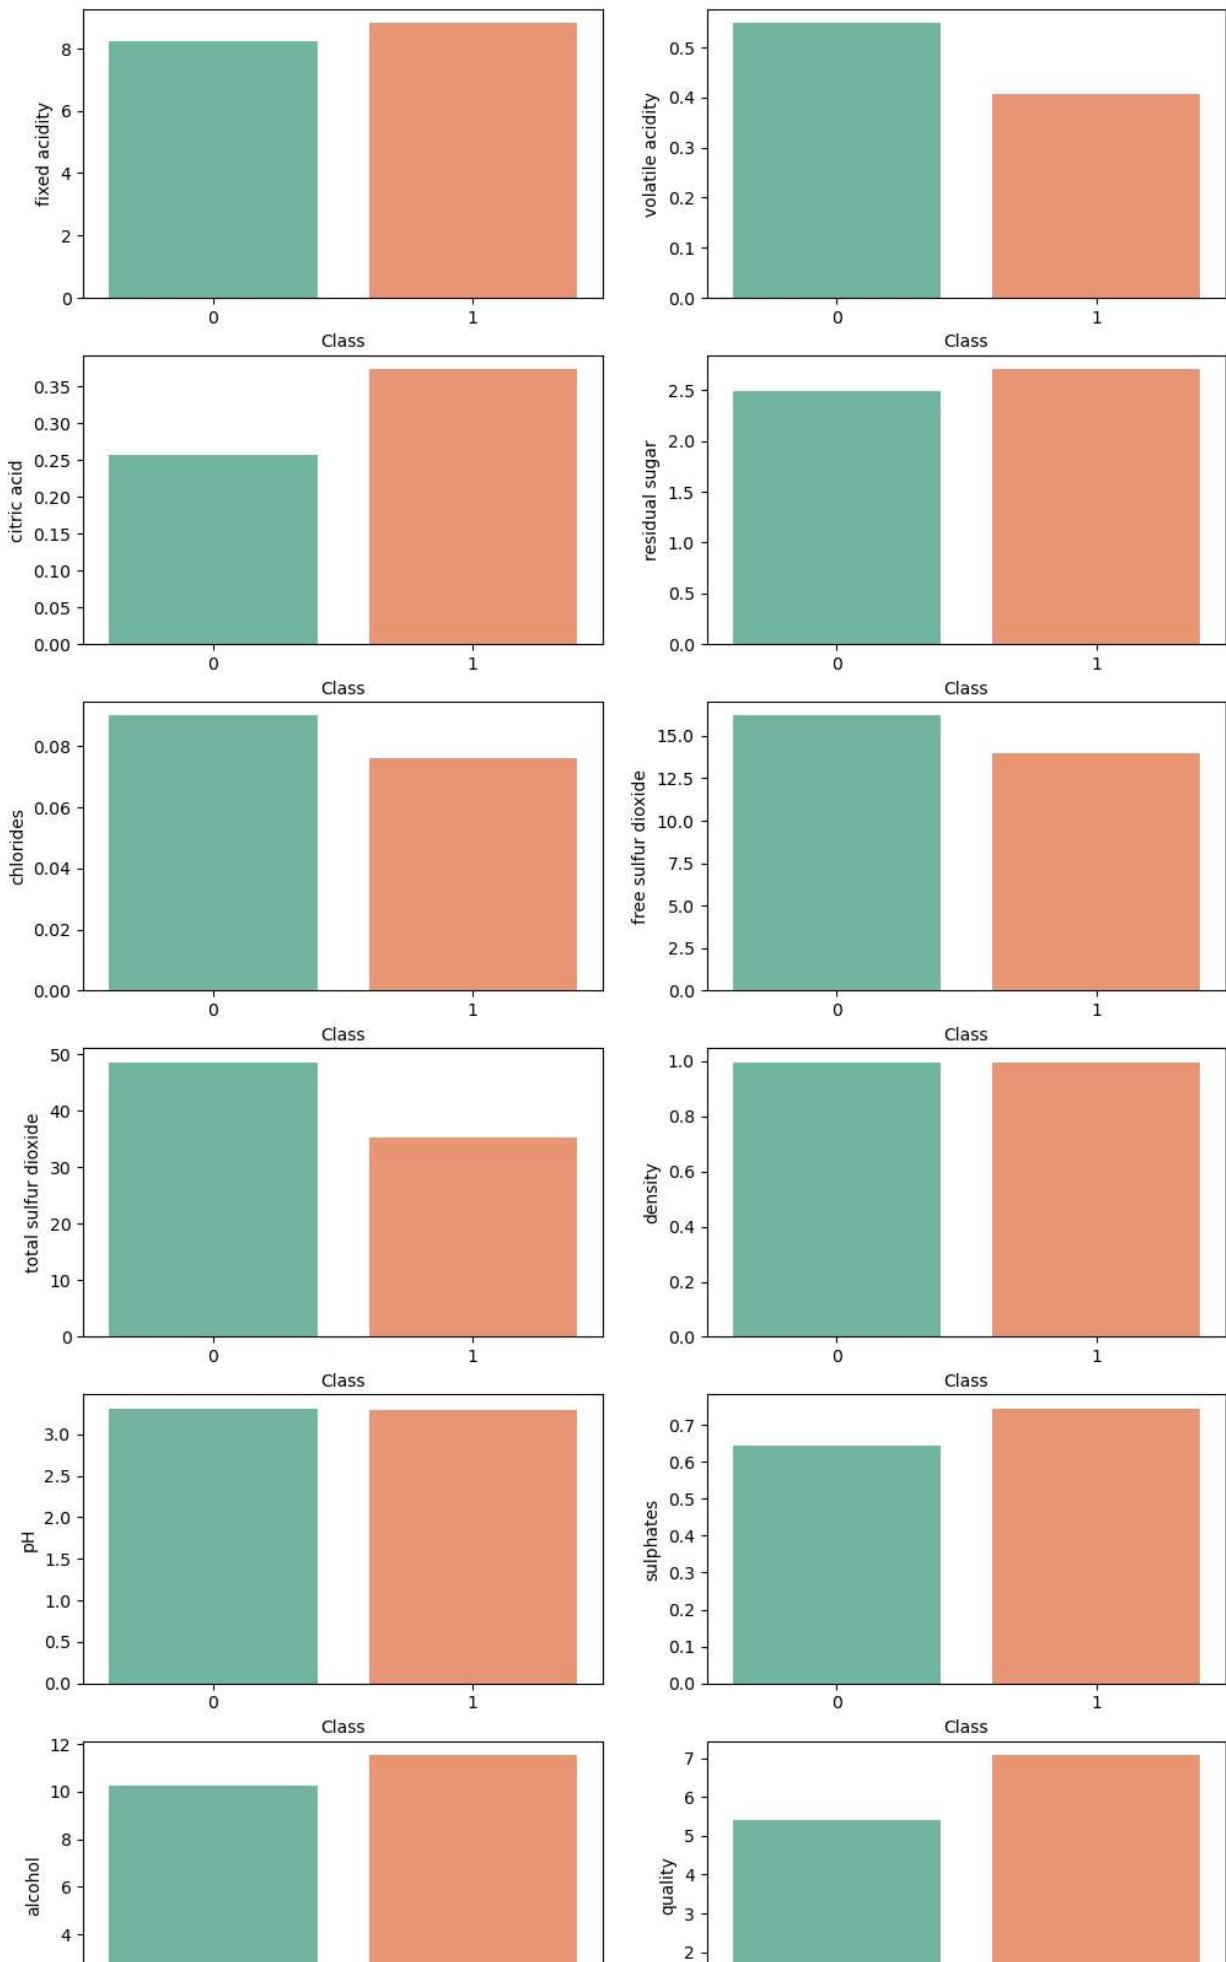
```
In [15]: cmp = df.groupby('Class').mean()
cmp
```

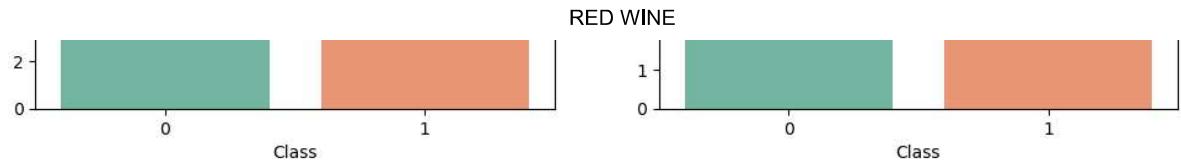
Out[15]:

Class	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sul
0	8.230809	0.548834	0.256587	2.495149	0.090014	16.194043	48.621702	0.996826	3.313106	0.6
1	8.820109	0.405870	0.372880	2.703804	0.076049	13.972826	35.358696	0.995963	3.288587	0.7

```
In [16]: plt.figure(figsize = (12,80))
plt_num = 1
for i in cmp.columns:
    plt.subplot(22,2,plt_num)
    sns.barplot(x = cmp.index, y = cmp[i], palette ='Set2')
    plt_num+=1
plt.show()
```

## RED WINE





```
In [17]: df['Class'].value_counts()
```

```
Out[17]: Class
0    1175
1    184
Name: count, dtype: int64
```

```
In [18]: good = df[df['Class']==1]
bad = df[df['Class']==0]

good = good.sample(bad.shape[0], replace=True)

print(bad.shape)
print(good.shape)
```

```
(1175, 13)
(1175, 13)
```

```
In [19]: data = pd.concat([bad, good])
data.shape
```

```
Out[19]: (2350, 13)
```

```
In [20]: df = df.drop('quality', axis = 1)
```

```
In [21]: data.head()
```

```
Out[21]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qua
<b>0</b>	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	
<b>1</b>	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	
<b>2</b>	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	
<b>3</b>	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	
<b>4</b>	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	

```
In [22]: x = data.drop('Class', axis = 1)
y = data['Class']
```

```
In [23]: x.head()
```

## RED WINE

Out[23]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qua
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	
4	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	

In [24]: `y.head()`

```
Out[24]: 0    0
1    0
2    0
3    0
4    0
Name: Class, dtype: int32
```

```
In [25]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_scaler = scaler.fit_transform(x)
```

```
In [26]: pd.DataFrame(x_scaler, columns = x.columns).head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sul
0	-0.622021	1.241914	-1.562028	-0.516303	-0.167332	-0.398335	-0.235469	0.651899	1.338337	-0.8
1	-0.410403	2.253807	-1.562028	-0.021953	0.344383	0.940503	0.736227	0.160717	-0.632870	-0.1
2	-0.410403	1.579211	-1.364058	-0.233817	0.204824	-0.015810	0.353438	0.258953	-0.251346	-0.2
3	1.388347	-1.119171	1.209557	-0.516303	-0.190591	0.175453	0.530110	0.750136	-0.887219	-0.1
4	-0.622021	1.017048	-1.562028	-0.586925	-0.190591	-0.207072	-0.058797	0.651899	1.338337	-0.8

```
In [27]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaler, y, test_size = 0.25, stra
```

```
In [28]: print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)
(1762, 12) (588, 12) (1762,) (588,)
```

```
In [29]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_train, y_train)
```

Out[29]:

```
    ▾ LogisticRegression
      LogisticRegression()
```

In [30]:

```
train_pred_lr = lr.predict(x_train)
test_pred_lr = lr.predict(x_test)
```

In [31]:

```
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
print(confusion_matrix(y_train, train_pred_lr))
print()
print(confusion_matrix(y_test, test_pred_lr))
```

```
[[881  0]
 [ 0 881]]
```

```
[[294  0]
 [ 0 294]]
```

In [32]:

```
print(classification_report(y_train, train_pred_lr))
print()
print(classification_report(y_test, test_pred_lr))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762

	precision	recall	f1-score	support
0	1.00	1.00	1.00	294
1	1.00	1.00	1.00	294
accuracy			1.00	588
macro avg	1.00	1.00	1.00	588
weighted avg	1.00	1.00	1.00	588

In [33]:

```
train_accu = []
test_accu = []
train_mean = []
test_mean = []
```

In [34]:

```
print(accuracy_score(y_train, train_pred_lr))
train_accu.append(accuracy_score(y_train, train_pred_lr))
print()
print(accuracy_score(y_test, test_pred_lr))
test_accu.append(accuracy_score(y_test, test_pred_lr))
```

1.0

1.0

In [35]:

```
from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(lr, x_train, y_train, cv=10)
```

```

test_accuracy = cross_val_score(lr, x_test, y_test, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())

```

Train Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Train Mean Accuracy 1.0

Train Max Accuracy 1.0

Test Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Test Mean Accuracy 1.0

Test Max Accuracy 1.0

In [36]:

```

from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(max_depth = 8,min_samples_leaf = 2)
dtree.fit(x_train,y_train)

```

Out[36]:

DecisionTreeClassifier(max\_depth=8, min\_samples\_leaf=2)

In [37]:

```

train_pred_dtree = dtree.predict(x_train)
test_pred_dtree = dtree.predict(x_test)

```

In [38]:

```

print(confusion_matrix(y_train,train_pred_dtree))
print()
print(confusion_matrix(y_test,test_pred_dtree))

```

[[881 0]  
 [ 0 881]]

[[294 0]  
 [ 0 294]]

In [39]:

```

print(classification_report(y_train,train_pred_dtree))
print()
print(classification_report(y_test,test_pred_dtree))

```

## RED WINE

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762
	precision	recall	f1-score	support
0	1.00	1.00	1.00	294
1	1.00	1.00	1.00	294
accuracy			1.00	588
macro avg	1.00	1.00	1.00	588
weighted avg	1.00	1.00	1.00	588

```
In [40]: print(accuracy_score(y_train,train_pred_dtreet))
train_accu.append(accuracy_score(y_train,train_pred_dtreet))
print()
print(accuracy_score(y_test,test_pred_dtreet))
test_accu.append(accuracy_score(y_test,test_pred_dtreet))
```

1.0

1.0

```
In [41]: from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(dtreet,x_train, y_train, cv=10)
test_accuracy = cross_val_score(dtreet,x_train, y_train, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())
```

Train Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Train Mean Accuracy 1.0

Train Max Accuracy 1.0

Test Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Test Mean Accuracy 1.0

Test Max Accuracy 1.0

```
In [42]: from sklearn.ensemble import RandomForestClassifier
rfr = RandomForestClassifier(n_estimators=200, max_depth=8, min_samples_leaf= 1, min_s
```

```
rfr.fit(x_train,y_train)
```

Out[42]:

```
RandomForestClassifier
RandomForestClassifier(max_depth=8, n_estimators=200, random_state=200)
```

In [43]:

```
train_pred_rfr = rfr.predict(x_train)
test_pred_rfr = rfr.predict(x_test)

print(confusion_matrix(y_train,train_pred_rfr))
print()
print(confusion_matrix(y_test,test_pred_rfr))
```

```
[[881  0]
 [ 0 881]]
```

```
[[294  0]
 [ 0 294]]
```

In [44]:

```
print(classification_report(y_train,train_pred_rfr))
print()
print(classification_report(y_test,test_pred_rfr))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762

	precision	recall	f1-score	support
0	1.00	1.00	1.00	294
1	1.00	1.00	1.00	294
accuracy			1.00	588
macro avg	1.00	1.00	1.00	588
weighted avg	1.00	1.00	1.00	588

In [45]:

```
print(accuracy_score(y_train,train_pred_rfr))
train_accu.append(accuracy_score(y_train,train_pred_rfr))
print()
print(accuracy_score(y_test,test_pred_rfr))
test_accu.append(accuracy_score(y_test,test_pred_rfr))
```

1.0

1.0

In [46]:

```
from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(rfr, x_train, y_train, cv=10)
test_accuracy = cross_val_score(rfr, x_test, y_test, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
```

```

print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())

```

Train Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Train Mean Accuracy 1.0

Train Max Accuracy 1.0

Test Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Test Mean Accuracy 1.0

Test Max Accuracy 1.0

In [47]: `!pip install xgboost`

```

Requirement already satisfied: xgboost in c:\users\prane\anaconda3\lib\site-packages
(2.0.2)
Requirement already satisfied: numpy in c:\users\prane\anaconda3\lib\site-packages (f
rom xgboost) (1.24.3)
Requirement already satisfied: scipy in c:\users\prane\anaconda3\lib\site-packages (f
rom xgboost) (1.11.1)

```

WARNING: There was an error checking the latest version of pip.

In [48]: `from xgboost import XGBClassifier  
xgb = XGBClassifier()  
xgb.fit(x_train,y_train)`

Out[48]:

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=N
one,
              enable_categorical=False, eval_metric=None, feature_types=N
one,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=N
one,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,

```

In [49]: `train_pred_xgb = xgb.predict(x_train)  
test_pred_xgb = xgb.predict(x_test)  
  
print(confusion_matrix(y_train,train_pred_xgb))  
print()  
print(confusion_matrix(y_test,test_pred_xgb))`

```
[[881  0]
 [ 0 881]]
```

```
[[294  0]
 [ 0 294]]
```

```
In [50]: print(classification_report(y_train,train_pred_xgb))
print()
print(classification_report(y_test,test_pred_xgb))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762
	precision	recall	f1-score	support
0	1.00	1.00	1.00	294
1	1.00	1.00	1.00	294
accuracy			1.00	588
macro avg	1.00	1.00	1.00	588
weighted avg	1.00	1.00	1.00	588

```
In [51]: print(accuracy_score(y_train,train_pred_xgb))
train_accu.append(accuracy_score(y_train,train_pred_xgb))
print()
print(accuracy_score(y_test,test_pred_xgb))
test_accu.append(accuracy_score(y_test,test_pred_xgb))
```

1.0

1.0

```
In [52]: from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(xgb, x_train, y_train, cv=10)
test_accuracy = cross_val_score(xgb, x_test, y_test, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())
```

```
Train Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```
Train Mean Accuracy 1.0
```

```
Train Max Accuracy 1.0
```

```
Test Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

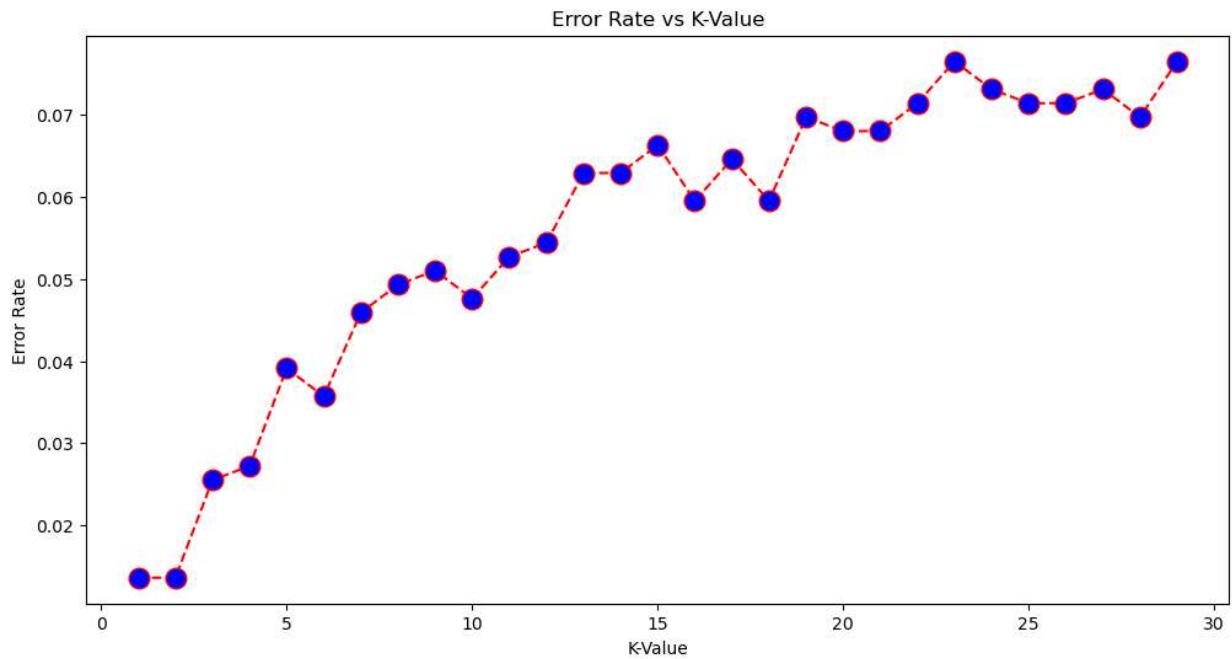
```
Test Mean Accuracy 1.0
```

```
Test Max Accuracy 1.0
```

```
In [53]: from sklearn.neighbors import KNeighborsClassifier
error_rate = []

for i in range(1,30):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train, y_train)
    y_pred = knn.predict(x_test)
    error_rate.append(np.mean(y_pred != y_test))
```

```
In [54]: plt.figure(figsize=(12,6))
plt.plot(range(1,30), error_rate, color='red', linestyle='dashed', marker='o',
         markersize=12, markerfacecolor='blue')
plt.title("Error Rate vs K-Value")
plt.xlabel("K-Value")
plt.ylabel("Error Rate")
plt.show()
```



```
In [55]: knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(x_train,y_train)
```

```
Out[55]: KNeighborsClassifier
```

```
KNeighborsClassifier(n_neighbors=1)
```

```
In [56]: train_pred_knn = knn.predict(x_train)
test_pred_knn = knn.predict(x_test)
print(confusion_matrix(y_train,train_pred_knn))
print()
print(confusion_matrix(y_test,test_pred_knn))
```

```
[[881  0]
 [  0 881]]
```

```
[[286   8]
 [  0 294]]
```

```
In [57]: print(classification_report(y_train,train_pred_knn))
print()
print(classification_report(y_test,test_pred_knn))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762

	precision	recall	f1-score	support
0	1.00	0.97	0.99	294
1	0.97	1.00	0.99	294
accuracy			0.99	588
macro avg	0.99	0.99	0.99	588
weighted avg	0.99	0.99	0.99	588

```
In [58]: print(accuracy_score(y_train,train_pred_knn))
print()
train_accu.append(accuracy_score(y_train,train_pred_knn))
print(accuracy_score(y_test,test_pred_knn))
test_accu.append(accuracy_score(y_test,test_pred_knn))
```

1.0

0.9863945578231292

```
In [59]: from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(knn, x_train, y_train, cv=10)
test_accuracy = cross_val_score(knn, x_test, y_test, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
```

```
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())

Train Accuracy [0.98870056 1.          0.96590909 0.98295455 0.98863636 0.98863636
0.99431818 0.98863636 0.98863636 0.98295455]

Train Mean Accuracy 0.9869382383153569

Train Max Accuracy 1.0

Test Accuracy [0.91525424 0.96610169 0.91525424 0.93220339 0.93220339 0.98305085
0.96610169 0.94915254 0.94827586 0.93103448]

Test Mean Accuracy 0.9438632378725892

Test Max Accuracy 0.9830508474576272
```

```
In [60]: # Support Vector Machine classification algorithm
from sklearn.svm import SVC
svm = SVC(kernel = 'rbf')
svm.fit(x_train,y_train)
```

```
Out[60]: ▾ SVC
SVC()
```

```
In [61]: train_pred_svm = svm.predict(x_train)
test_pred_svm = svm.predict(x_test)
print(confusion_matrix(y_train,train_pred_svm))
print()
print(confusion_matrix(y_test,test_pred_svm))
```

```
[[881  0]
 [ 0 881]]
```

```
[[294  0]
 [ 0 294]]
```

```
In [62]: print(classification_report(y_train,train_pred_svm))
print()
print(classification_report(y_test,test_pred_svm))
```

## RED WINE

	precision	recall	f1-score	support
0	1.00	1.00	1.00	881
1	1.00	1.00	1.00	881
accuracy			1.00	1762
macro avg	1.00	1.00	1.00	1762
weighted avg	1.00	1.00	1.00	1762
	precision	recall	f1-score	support
0	1.00	1.00	1.00	294
1	1.00	1.00	1.00	294
accuracy			1.00	588
macro avg	1.00	1.00	1.00	588
weighted avg	1.00	1.00	1.00	588

```
In [63]: print(accuracy_score(y_train,train_pred_svm))
train_accu.append(accuracy_score(y_train,train_pred_svm))
print()
print(accuracy_score(y_test,test_pred_svm))
test_accu.append(accuracy_score(y_test,test_pred_svm))
```

1.0

1.0

```
In [64]: from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(svm, x_train, y_train, cv=10)
test_accuracy = cross_val_score(svm, x_test, y_test, cv=10)
print("Train Accuracy", training_accuracy)
print()
print("Train Mean Accuracy", training_accuracy.mean())
print()
print("Train Max Accuracy", training_accuracy.max())
print()
print("Test Accuracy", test_accuracy)
print()
print("Test Mean Accuracy", test_accuracy.mean())
print()
print("Test Max Accuracy", test_accuracy.max())
train_mean.append(training_accuracy.mean())
test_mean.append(test_accuracy.mean())
```

Train Accuracy [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Train Mean Accuracy 1.0

Train Max Accuracy 1.0

Test Accuracy [0.98305085 1. 1. 1. 0.98305085 0.98275862 1. ]

Test Mean Accuracy 0.9931911163062537

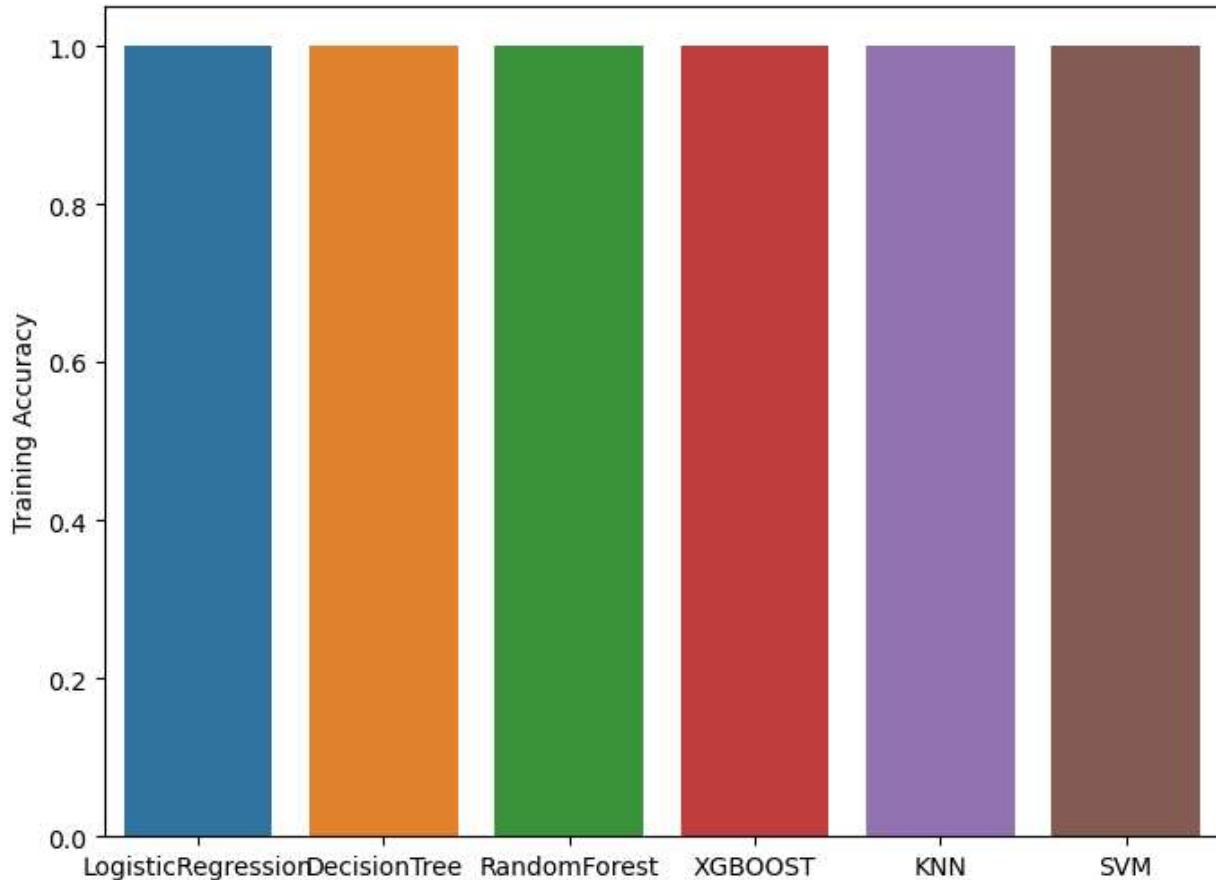
Test Max Accuracy 1.0

```
In [65]: algorithm = ['LogisticRegression', 'DecisionTree', 'RandomForest', 'XGBOOST', 'KNN', 'SVM']
accu_data = {'Training Accuracy':train_accu,'Test Accuracy':test_accu, 'Train Mean':train_mean, 'Test Mean':test_mean}
model = pd.DataFrame(accu_data, index = algorithm)
```

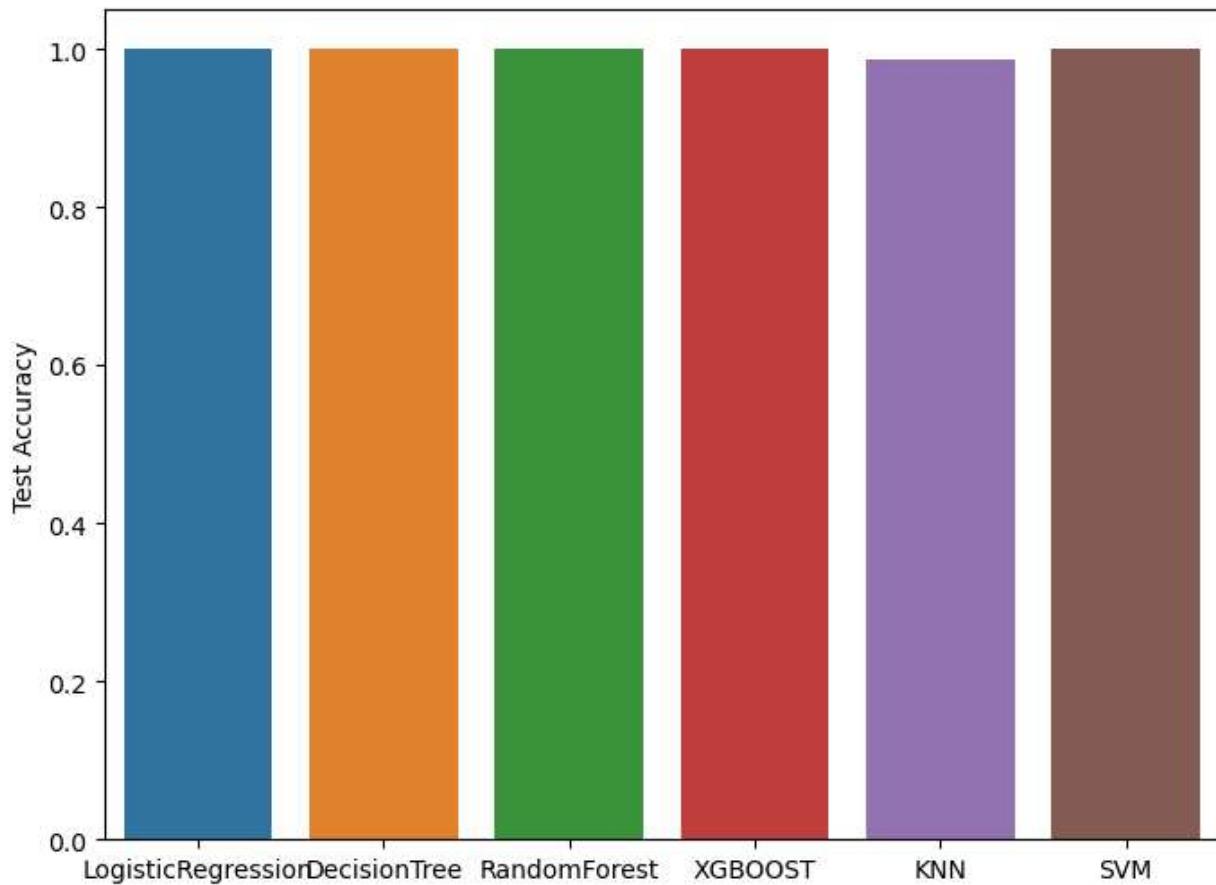
Out[65]:

	Training Accuracy	Test Accuracy	Train Mean	Test Mean
<b>LogisticRegression</b>	1.0	1.000000	1.000000	1.000000
<b>DecisionTree</b>	1.0	1.000000	1.000000	1.000000
<b>RandomForest</b>	1.0	1.000000	1.000000	1.000000
<b>XGBOOST</b>	1.0	1.000000	1.000000	1.000000
<b>KNN</b>	1.0	0.986395	0.986938	0.943863
<b>SVM</b>	1.0	1.000000	1.000000	0.993191

```
In [66]: plt.figure(figsize=(8,6))
sns.barplot(x=model.index, y=model['Training Accuracy'])
plt.show()
```



```
In [67]: plt.figure(figsize=(8,6))
sns.barplot(x=model.index, y=model['Test Accuracy'])
plt.show()
```



In [ ]: